



# Lifelong Learning ( $L^2$ )


Shirin Dora

October 30, 2023




# Outline


- Motivation: Why lifelong learning?
  - Classical machine learning
  - Catastrophic forgetting
- Definitions
  - Lifelong Learning
  - Three  $L^2$  settings
- Current  $L^2$  approaches
  - Replay Methods
  - Regularization Methods
  - Architectural Methods
- Evaluation of  $L^2$  methods
- My Research interest




- **Classical machine learning:** Isolated single-task learning.
- What is a task?



- **Classical machine learning:** Isolated single-task learning.
- What is a task?
  - Classification: cars vs truck, MNIST.
  - Regression: stock prices, rainfall.



- **Classical machine learning:** Isolated single-task learning.
- What is a task?
  - Classification: cars vs truck, MNIST.
  - Regression: stock prices, rainfall.
- Each model performs a single task.
  - Support vector machines, Deep neural networks.




- **Classical machine learning:** Isolated single-task learning.
- What is a task?
  - Classification: cars vs truck, MNIST.
  - Regression: stock prices, rainfall.
- Each model performs a single task.
  - Support vector machines, Deep neural networks.
  - Very effective but there are limitations!



## Key issues in classical machine learning


- No lifelong/continual learning
  - Learning occurs in isolation.
  - Knowledge accumulation is not possible.
  - Knowledge transfer is not possible.



## Key issues in classical machine learning

- No lifelong/continual learning
  - Learning occurs in isolation.
  - Knowledge accumulation is not possible.
  - Knowledge transfer is not possible.
- Assumes stationarity: Can't handle environment changes.





## Key issues in classical machine learning

- No lifelong/continual learning
  - Learning occurs in isolation.
  - Knowledge accumulation is not possible.
  - Knowledge transfer is not possible.
- Assumes stationarity: Can't handle environment changes.
- No more learning is possible after deployment.



## Easy Solution

- Train a single model on multiple tasks one after another (in a sequence).



## Easy Solution

- Train a single model on multiple tasks one after another (in a sequence).
- Learning requires updating parameters of the **shared** model for each task (like weights in a neural network).



## Easy Solution


- Train a single model on multiple tasks one after another (in a sequence).
- Learning requires updating parameters of the **shared** model for each task (like weights in a neural network).
  - Knowledge acquired from previous tasks will be overwritten.
  - Performance on previous tasks deteriorate as we learn new tasks - *Catastrophic forgetting*.



# Catastrophic Forgetting



- **Task 1**
  - 2 classes: blue and green squares.






# Catastrophic Forgetting



- **Task 1**
  - 2 classes: blue and green squares.
- Find a line that separates two classes.
  - $y = mx + c$
  - Model = Finding  $m$  and  $c$ .






# Catastrophic Forgetting



- Apply the easy solution
  - Continue training the model on Task 2.
- **Task 2**
  - 2 classes: orange and magenta squares.
  - Find a line that separates two classes.






# Catastrophic Forgetting



- Apply the easy solution
  - Continue training the model on Task 2.
- **Task 2**
  - 2 classes: orange and magenta squares.
  - Find a line that separates two classes.
- Red line is our model now.








# Catastrophic Forgetting



- We wanted a **single shared model** for all tasks.
- New model (line) doesn't work for Task 1.
- Learning Task 2 made model forget Task 1.
  - Catastrophic forgetting.
- Catastrophic forgetting is the most fundamental problem in  $L^2$ .





# Definitions



# Lifelong Learning ( $L^2$ )

- Learn a sequence of tasks,  $T_1, T_2, \dots, T_N, \dots$  sequentially.
- Each task  $t$  has a dataset for training,  $\mathcal{D}_t = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$



# Lifelong Learning ( $L^2$ )

- Learn a sequence of tasks,  $T_1, T_2, \dots, T_N, \dots$  sequentially.
- Each task  $t$  has a dataset for training,  $\mathcal{D}_t = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$
- **Goal:** Learn each new task  $T_{N+1}$  sequentially:
  - *Without catastrophic forgetting:* Learning of new task  $T_{N+1}$  should not result in degradation of accuracy for the previous  $N$  tasks.



# Lifelong Learning ( $L^2$ )

- Learn a sequence of tasks,  $T_1, T_2, \dots, T_N, \dots$  sequentially.
- Each task  $t$  has a dataset for training,  $\mathcal{D}_t = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$
- **Goal:** Learn each new task  $T_{N+1}$  sequentially:
  - *Without catastrophic forgetting*: Learning of new task  $T_{N+1}$  should not result in degradation of accuracy for the previous  $N$  tasks.
  - *With knowledge transfer*: Utilize knowledge acquired from previous tasks to learn the new task  $T_{N+1}$  more effectively.



# Lifelong Learning ( $L^2$ )

- Learn a sequence of tasks,  $T_1, T_2, \dots, T_N, \dots$  sequentially.
- Each task  $t$  has a dataset for training,  $\mathcal{D}_t = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$
- **Goal:** Learn each new task  $T_{N+1}$  sequentially:
  - *Without catastrophic forgetting*: Learning of new task  $T_{N+1}$  should not result in degradation of accuracy for the previous  $N$  tasks.
  - *With knowledge transfer*: Utilize knowledge acquired from previous tasks to learn the new task  $T_{N+1}$  more effectively.
- **Assumption:** Once a task is learned, its data is no longer accessible (at least majority of it).



# Lifelong Learning ( $L^2$ )

## Three $L^2$ settings [Van de Ven, 2022]

1. Task incremental learning
  - Models are *partially* shared across tasks.
  - Task identity is required during testing.



# Lifelong Learning ( $L^2$ )

## Three $L^2$ settings [Van de Ven, 2022]

1. Task incremental learning
  - Models are *partially* shared across tasks.
  - Task identity is required during testing.
2. Domain incremental learning
  - All tasks have the same set of classes.
  - Task identity is not required during testing.
  - E.g. Classifying objects under different lightning conditions.






# Lifelong Learning ( $L^2$ )

## Three $L^2$ settings [Van de Ven, 2022]

1. Task incremental learning
  - Models are *partially* shared across tasks.
  - Task identity is required during testing.
2. Domain incremental learning
  - All tasks have the same set of classes.
  - Task identity is not required during testing.
  - E.g. Classifying objects under different lightning conditions.
3. Class incremental learning
  - Produce a single model from all tasks.
  - All classes in all tasks are handled by one model.



# Three $L^2$ Settings



Task-incremental learning (Task-IL)

Choice between 2 known digits (e.g., '0' or '1')

Class-incremental learning (Class-IL)

Choice between all digits seen so far



# Three $L^2$ Settings



Task-incremental learning (Task-IL)  
Choice between 2 known digits (e.g., '0' or '1'?)

Class-incremental learning (Class-IL)  
Choice between all digits seen so far



## Domain Incremental Learning




# Current Approaches



# Current Approaches

- **Replay/Rehearsal-based methods**
  - Replay samples from previous tasks while learning a new task.
- **Regularization-based methods**
  - Update to model parameters is regularized according to a parameter's importance for previous tasks.
- **Architectural methods**
  - Expand network architecture to learn new tasks.




## Task 1 Training

- Save subset of samples in replay memory.




# Replay Methods



## Task 1 Training

- Save subset of samples in replay memory.



## Task 2 Training

- Replay stored samples from previous tasks.
- Save subset of samples in replay memory.



## Why replay works?

- For training, mix data from new task and stored samples.
  - **Stored samples:** Won't allow changes to weights important for previous tasks.
  - **New task samples:** Learn the new task.
- The quality of stored samples is very important!





## Why replay works?


- For training, mix data from new task and stored samples.
  - **Stored samples:** Won't allow changes to weights important for previous tasks.
  - **New task samples:** Learn the new task.
- The quality of stored samples is very important!

## Is there another way to obtain stored samples?

- Train a generative model in-parallel to your classifier.
- Sample replay samples from the generative model [Shin, 2017].



# Replay Methods



- Generative replay
  - Samples for replay are obtained from generative models.
- Exact replay
  - Actual samples from the original dataset are used for replay.



- Learning in deep neural networks and brains with similarity-weighted interleaved learning, 2022 - Only previous memories with high similarity to the new data.

## Coresets

- Gradient based sample selection for online continual learning, 2019 - Select samples that exhibit maximum variance in gradient.
- Online Continual Learning with Maximally Interfered Retrieval, 2019. - Samples most affected by a given gradient update.



# Regularization Methods



- Use an additional penalty term during learning.
- Without regularization

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E}(y_{output}, y_{desired}) \quad (1)$$

- $\mathcal{E}$  could be MSE, cross entropy, etc.



- Use an additional penalty term during learning.
- Without regularization

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E}(y_{\text{output}}, y_{\text{desired}}) \quad (1)$$

–  $\mathcal{E}$  could be MSE, cross entropy, etc.

- With regularization,

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \text{parameter-wise penalty} \quad (2)$$



## Regularization: Old wine in new bottles!

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda \sum_i \theta_i^2 \quad (3)$$

- **L2 Regularization**
- $\theta_i$  denotes parameters of your network.
- $\lambda$  determines how strongly regularization is applied.
- *Almost* a standard technique to prevent overfitting.



## Regularization: Old wine in new bottles!

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda \sum_i \theta_i^2 \quad (3)$$

- **L2 Regularization**
- $\theta_i$  denotes parameters of your network.
- $\lambda$  determines how strongly regularization is applied.
- *Almost* a standard technique to prevent overfitting.
- Keep  $\theta_i$  close to zero.

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda \sum_i (\theta_i - 0)^2 \quad (4)$$





$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda \sum_i \theta_i^2 \quad (5)$$

- $\lambda$  is same for all  $\theta_i$ .
- All parameters might not be equally important!



$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda \sum_i \theta_i^2 \quad (5)$$

- $\lambda$  is same for all  $\theta_i$ .
- All parameters might not be equally important!

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda_i \sum_i \theta_i^2 \quad (6)$$

- High  $\lambda_i$  for less important  $\theta_i$ .
- Low  $\lambda_i$  for highly important  $\theta_i$ .



$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda \sum_i \theta_i^2 \quad (5)$$


- $\lambda$  is same for all  $\theta_i$ .
- All parameters might not be equally important!

$$\mathcal{L}_B = \sum_{x \in B} \mathcal{E} + \lambda_i \sum_i \theta_i^2 \quad (6)$$

- High  $\lambda_i$  for less important  $\theta_i$ .
- Low  $\lambda_i$  for highly important  $\theta_i$ .
- How do we set  $\lambda_i$ ?




# Regularization Methods






# Regularization Methods






# Regularization Methods



**The usual gradient descent!**




# Regularization Methods



## Low gradient

- Little change in weight.
- Less importance.




## High gradient

- Large change in weight.
- High importance.



## Continual learning

- Use parameters less important for previous tasks to learn new task.
  - Low gradient for previous tasks.
  - High gradient for current task.








## Continual learning

- Use parameters less important for previous tasks to learn new task.
  - Low gradient for previous tasks.
  - High gradient for current task.
- Double derivative (gradient of the gradient) represents speed of change in the gradient.





$$\mathcal{L}_{new} = \sum_{x \in new} \mathcal{E} + c\lambda_i \sum_i (\theta_i^2 - \theta_{prev}) \quad (7)$$

- $\theta_{prev}$  represents network parameters after learning previous tasks.
- $\lambda_i$  would be estimated using double derivative [Kirkpatrick, 2016].

$$\lambda_i \propto \frac{\delta^2 \mathcal{L}}{\delta \theta_i^2}$$

- $(\theta_i - \theta_{prev})$  has the effect of keeping  $\theta_i$  close to  $\theta_{prev}$  when  $\lambda_i$  is high.
- $c$  controls strength of regularization.



- Continual learning through synaptic intelligence -  
Computationally less intensive method for estimating  $\lambda$ .
- Memory Aware Synapses: Learning what (not) to forget -  
Effect of a parameter change on the network output.
- Optimization and Generalization of Regularization-Based  
Continual Learning: a Loss Approximation Viewpoint -  
Unifies various parameter-level regularization methods.
- Continual learning with node-importance based adaptive  
group sparse regularization, 2020 - Regularization at the  
level of neurons.




# Architectural Methods



# Architectural Methods

- Expand the network to learn new tasks.
- Brute force approach to  $L^2$ .
- Leads to parameter explosion.



Progressive Neural  
Networks[Rusu, 2016]



- **Progress & compress:** A scalable framework for continual learning. - Introduced a compression step for progressive neural networks.

## **Prune/Add weights to the network**

- Lifelong learning with dynamically expandable networks.
- Compacting, picking and growing for unforgetting continual learning.
- Learn to grow: A continual structure learning framework for overcoming catastrophic forgetting.




# Evaluating and Debugging $L^2$ Methods



# Interpreting $L^2$ Results



- What do you notice?








# Interpreting $L^2$ Results

- What do you notice?
- Learned task performance is getting worse after each task is trained.
- This is a sign of catastrophic forgetting!






# Interpreting $L^2$ Results



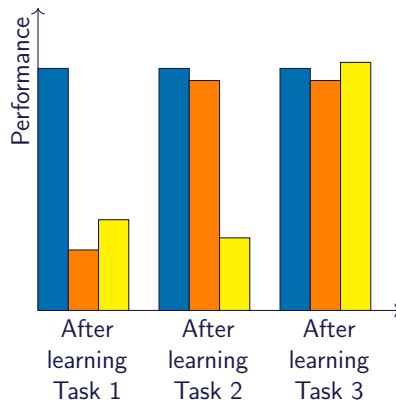
- What do you notice?





# Interpreting $L^2$ Results


- What do you notice?
- Learned task performance is maintained after learning a new task.
- This indicates no forgetting!





# Interpreting $L^2$ Results

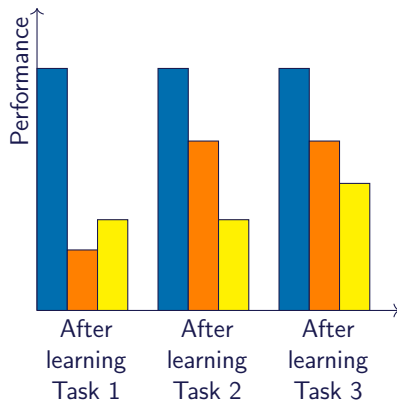
- What do you notice?





# Interpreting $L^2$ Results


- What do you notice?
- Learned task performance is maintained but new tasks are not learned well enough.
- Problematic learning
  - Over regularization.
  - Limited network capacity.





# Interpreting $L^2$ Results


- What do you notice?





# Interpreting $L^2$ Results


- What do you notice?
- Learned task performance is maintained
- New task performance is higher than previous task.
- **Forward transfer**
  - Future tasks utilize knowledge from previous tasks.





# Interpreting $L^2$ Results

- What do you notice?








# Interpreting $L^2$ Results

- What do you notice?
- Learned task performance is maintained
- New task performance is higher than previous task.
- Forward and **backward transfer**
  - Future tasks utilize knowledge from previous tasks.
  - Previous tasks gain from future tasks.





- Metrics allow us to determine which scenario our model results represent.
- All  $L^2$  metrics are computed from the accuracy matrix [Lopez-paz, 2017](#).



# Accuracy Matrix



		Testing task				
		$T_1$	$T_2$	$T_3$	$T_4$	$\dots$
Tasks trained so far	$T_1$	$R_{1,1}$	$R_{1,2}$	$R_{1,3}$	$R_{1,4}$	
	$T_2$	$R_{2,1}$	$R_{2,2}$	$R_{2,3}$	$R_{2,4}$	
	$T_3$	$R_{3,1}$	$R_{3,2}$	$R_{3,3}$	$R_{3,4}$	
	$T_4$	$R_{4,1}$	$R_{4,2}$	$R_{4,3}$	$R_{4,4}$	
	$\dots$					


$R_{m,n}$ : The performance of the model on task  $T_n$ , after continually training till task  $T_m$



# Accuracy Matrix



- Increasing  $\rightarrow$  Positive Forward transfer






# Accuracy Matrix

- Increasing  $\rightarrow$  Positive Forward transfer
- Measuring forward transfer (FWT)

$$\frac{1}{T-1} \sum_{i=2}^T (R_{i-1,i} - b_i) \quad (8)$$


- $b_i$  is accuracy before training.



# Accuracy Matrix

- Measuring backward transfer (BWT)

$$\frac{1}{T-1} \sum_{i=1}^{T-1} (R_{T,i} - R_{i,i}) \quad (9)$$




# Accuracy Matrix

- Measuring backward transfer (BWT)

$$\frac{1}{T-1} \sum_{i=1}^{T-1} (R_{T,i} - R_{i,i}) \quad (9)$$

- $R_{i,i}$ : accuracy after training a task.
- $R_{T,i}$ : accuracy after training last task.





# My Research Interests






# Efficient Generative $L^2$




- Deep Bidirectional Predictive Coding (DBPC) [Senhui 2023].
- Propagation can occur in both directions in DBPC.





# Efficient Generative $L^2$

- Deep Bidirectional Predictive Coding (DBPC) [Senhui 2023].
- Propagation can occur in both directions in DBPC.
- Single network can do classification and generation.
  - Replay without additional memory!





# Efficient Generative $L^2$



Task 1: 0 and 1



Task 2: 2 and 3



Task 3: 4 and 5



Task 4: 6 and 7




Task 5: 8 and 9




# Distributed $L^2$

- Group of nodes perform  $L^2$  collaboratively.
- Nodes encounter tasks in different sequence.



# Distributed $L^2$

- Group of nodes perform  $L^2$  collaboratively.
- Nodes encounter tasks in different sequence.
- Sharing knowledge improves speed of learning  
[Nath 2023].
  - 12 agents learned 16 tasks faster than a smaller group.
  - Agents exchanged knowledge through task-specific masks.






# Distributed $L^2$

- Mask represents a subnetwork in a randomly initialized network

[Wortsman 2020].

- Training involves estimating a mask for a given task.
- Testing involves processing a sample using the estimated mask.





# References



- van de Ven, G. M., Tuytelaars, T., & Tolias, A. S. (2022). Three types of incremental learning. *Nature Machine Intelligence*, 4(12), 1185-1197.
- Kirkpatrick, J., Pascanu, R., Rabinowitz, N., Veness, J., Desjardins, G., Rusu, A. A., ... & Hadsell, R. (2017). Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13), 3521-3526.
- Shin, H., Lee, J. K., Kim, J., & Kim, J. (2017). Continual learning with deep generative replay. *Advances in neural information processing systems*, 30.
- Rusu, A. A., Rabinowitz, N. C., Desjardins, G., Soyer, H., Kirkpatrick, J., Kavukcuoglu, K., ... & Hadsell, R. (2016). Progressive neural networks. *arXiv preprint arXiv:1606.04671*.
- Lopez-Paz, D., & Ranzato, M. A. (2017). Gradient episodic memory for continual learning. *Advances in neural information processing systems*, 30.



# References



- Qiu, S., Bhattacharyya, S., Coyle, D., & Dora, S. (2023). Deep Predictive Coding with Bi-directional Propagation for Classification and Reconstruction. arXiv preprint arXiv:2305.18472.
- Nath, S., Peridis, C., Ben-Iwhiwhu, E., Liu, X., Dora, S., Liu, C., ... & Soltoggio, A. (2023). Sharing Lifelong Reinforcement Learning Knowledge via Modulating Masks. arXiv preprint arXiv:2305.10997.
- Wortsman, M., Ramanujan, V., Liu, R., Kembhavi, A., Rastegari, M., Yosinski, J., & Farhadi, A. (2020). Supermasks in superposition. Advances in Neural Information Processing Systems, 33, 15173-15184.